

1 Text Mining Bibliographic Corpora

1.1 Learning objectives

After completing this chapter, you will be able to:

- Build a text corpus from OpenAlex titles and abstracts
- Apply standard preprocessing steps: tokenisation, stopword removal, stemming
- Construct a document-term matrix (DTM) and a document-feature matrix (DFM)
- Compute TF-IDF weights to identify distinctive terms
- Visualise term frequency patterns across time or groups

1.2 Setup

```
library(tidyverse)
library(openalexR)
library(quanteda)
library(quanteda.textstats)
library(glue)
library(gt)

set.seed(20260509)

source(here::here("R", "api_helpers.R"))
source(here::here("R", "utils.R"))
source(here::here("R", "sci_palette.R"))
```

1.3 Conceptual background

Bibliometric metadata — titles, abstracts, and keywords — constitutes a rich but noisy text corpus. Text mining transforms this unstructured data into structured representations suitable for statistical analysis. The pipeline typically follows three stages: preprocessing, representation, and analysis.

Preprocessing converts raw text into a standardised form. Tokenisation splits text into individual words or n-grams. Lowercasing removes case variation. Stopword removal eliminates high-frequency function words (“the”, “of”, “and”) that carry little topical information. Stemming or lemmatisation reduces words to their root forms (“computing”, “computed”, “computation” → “comput” or “compute”). Each step involves trade-offs: aggressive stemming can merge distinct concepts, while minimal preprocessing retains noise.

The **document-term matrix** (DTM) or **document-feature matrix** (DFM) is the fundamental representation. Rows are documents; columns are terms; cells contain counts or weights. Raw term frequencies overweight common words. **TF-IDF** (term frequency–inverse document frequency) addresses this by upweighting terms that are frequent in a document but rare across the corpus, highlighting words that distinguish one document from others.

`quanteda` [Aria2017] provides a fast, well-designed toolkit for text analysis in R. Its `corpus` → `tokens` → `dfm` pipeline integrates naturally with the tidyverse. For very large corpora, `quanteda` uses sparse matrix representations that scale to millions of documents.

Text mining complements network-based methods (??). Co-word networks reveal term associations; text mining provides the frequency distributions, temporal trends, and discriminative features that characterise a field’s vocabulary.

1.4 Worked example

1.4.1 Building a text corpus

```
works <- oa_fetch(
  entity = "works",
  primary_location.source.id = "S148561398",
  from_publication_date = "2019-01-01",
  to_publication_date = "2023-12-31",
  options = list(sample = 500, seed = 42)
)

text_df <- works |>
  filter(!is.na(abstract), nchar(abstract) > 50) |>
  transmute(
    doc_id = id,
    text = paste(display_name, abstract, sep = ". "),
    year = year(publication_date)
  )

cat(glue("Documents with abstracts: {nrow(text_df)}\n"))
```

```
#> Documents with abstracts: 119
```

```
corp <- corpus(text_df, docid_field = "doc_id", text_field = "text")
docvars(corp, "year") <- text_df$year

cat(glue("Corpus size: {ndoc(corp)} documents\n"))
```

```
#> Corpus size: 119 documents
```

```
cat(glue("Total tokens: {sum(ntoken(corp))}\n"))
```

```
#> Total tokens: 29606
```

1.4.2 Tokenisation and preprocessing

```
toks <- tokens(corp,
  remove_punct = TRUE,
  remove_numbers = TRUE,
  remove_symbols = TRUE) |>
  tokens_tolower() |>
  tokens_remove(stopwords("en")) |>
  tokens_remove(c("also", "however", "using", "based", "study",
    "paper", "results", "research", "analysis")) |>
  tokens_wordstem()

cat(glue("Tokens after preprocessing: {sum(ntoken(toks))}\n"))
```

```
#> Tokens after preprocessing: 14958
```

1.4.3 Document-feature matrix

```
dfmat <- dfm(toks) |>
  dfm_trim(min_termfreq = 5, min_docfreq = 3)

cat(glue("DFM dimensions: {nrow(dfmat)} docs x {ncol(dfmat)} features\n"))
```

```
#> DFM dimensions: 119 docs x 659 features
```

```
cat(glue("Sparsity: {scales::percent(sparsity(dfmat))}\n"))
```

```
#> Sparsity: 91%
```

```
top_terms <- topfeatures(dfmat, 20) |>
  enframe(name = "term", value = "frequency") |>
  mutate(term = fct_reorder(term, frequency))

ggplot(top_terms, aes(x = frequency, y = term)) +
  geom_col(fill = palette_sci(1)) +
  labs(x = "Frequency", y = NULL) +
  theme_sci()
```

1.4.4 TF-IDF weighting

```
dfmat_tfidf <- dfm_tfidf(dfmat)

tfidf_by_year <- text_df |>
  group_by(year) |>
  group_keys() |>
  pull(year) |>
  map_dfr(function(yr) {
    docs <- docvars(dfmat_tfidf, "year") == yr
    if (sum(docs) < 5) return(tibble())
    top <- topfeatures(dfmat_tfidf[docs, ], 10)
    tibble(year = yr, term = names(top), tfidf = unname(top))
  })

tfidf_by_year |>
  group_by(year) |>
  slice_max(tfidf, n = 5) |>
  gt() |>
  fmt_number(columns = tfidf, decimals = 1)
```

1.4.5 Term frequency over time




Figure 1: Top 20 most frequent terms in the Scientometrics corpus.

```

target_terms <- c("bibliometr", "open", "network", "impact", "collabor")

freq_by_year <- map_dfr(unique(text_df$year), function(yr) {
  docs <- docvars(dfmat, "year") == yr
  if (sum(docs) < 5) return(tibble())
  freq <- colSums(dfmat[docs, ]) / sum(dfmat[docs, ])
  tibble(
    year = yr,
    term = target_terms,
    rel_freq = freq[target_terms]
  )
})

ggplot(freq_by_year, aes(x = year, y = rel_freq, colour = term)) +
  geom_line(linewidth = 1) +
  geom_point(size = 2) +
  scale_colour_manual(values = palette_sci(length(target_terms))) +
  scale_y_continuous(labels = scales::percent_format(accuracy = 0.01)) +
  labs(x = "Year", y = "Relative frequency", colour = "Term") +
  theme_sci()

```




Figure 2: Relative frequency of selected terms over time.

1.5 Diagnostics and interpretation

- **Vocabulary size vs. documents:** A typical DFM has far more features than documents. Trim aggressively (minimum document frequency of 3–5) to reduce noise and computation.

- **Sparsity:** Bibliometric DFMs are typically 95–99% sparse. This is normal. Sparse matrix representations keep memory usage manageable.
- **Preprocessing sensitivity:** Results change with preprocessing choices. Always report stopword list, stemming method, and trimming thresholds.
- **Abstract availability:** Not all OpenAlex records have abstracts. Report the coverage rate and consider whether missing abstracts introduce bias (e.g., older papers, certain publishers).

1.6 Limitations and responsible use

1.7 Limitations and responsible use

- **Abstracts are summaries.** They capture the main claims but miss nuance, methodology details, and negative results. Full-text analysis (??) provides richer data but is harder to obtain.
- **Language bias.** Text mining tools work best for English. Non-English abstracts may be poorly tokenised, incorrectly stemmed, or excluded entirely. This biases results toward Anglophone research [visser2021].
- **Stemming conflates concepts.** “Stem” and “stemming” are merged, but so are “cell” (biology) and “cell” (spreadsheet). Lemmatisation is more precise but slower.
- **Bag-of-words ignores context.** TF-IDF and frequency-based methods treat documents as unordered collections of words. Phrase meaning (“machine learning” vs. “learning machine”) is lost unless you use n-grams or embeddings (??).

1.8 Common pitfalls

1.9 Common pitfalls

- **Not removing domain-specific stop words.** Generic stopword lists miss terms like “study”, “results”, “paper” that dominate academic text without conveying topical information.
- **Applying TF-IDF before trimming.** Rare terms get extreme TF-IDF scores. Trim the DFM first, then apply TF-IDF.
- **Comparing raw frequencies across groups of different sizes.** A year with 200 papers will have higher raw counts than a year with 50. Always normalise to relative frequency.
- **Stemming titles but not abstracts (or vice versa).** Apply identical preprocessing to all text fields.

1.10 Exercises

1. **Bigram analysis.** Build a DFM using bigrams (two-word sequences) instead of unigrams. What meaningful phrases emerge that single words miss?
2. **Lexical diversity.** Compute the type-token ratio (TTR) for each publication year. Is the vocabulary growing more diverse or more standardised over time?
3. **Keyness analysis.** Use `quanteda.textstats::textstat_keyness()` to identify terms that distinguish one year from all others. What terms characterise 2023 specifically?
4. **Preprocessing sensitivity.** Run the same analysis with and without stemming. How do the top terms differ? Which version is more interpretable?

1.11 Solutions

Solutions are provided in 1.11.

1.12 Further reading

- @silge2017 — *Text Mining with R*; comprehensive guide to tidy text analysis.
- @aria2017 — `bibliometrix` text-mining features for bibliometric data.
- @priem2022 — OpenAlex abstract availability and coverage.

1.13 Session info

```
sessionInfo()
```

```
#> R version 4.4.1 (2024-06-14)
#> Platform: x86_64-pc-linux-gnu
#> Running under: Ubuntu 24.04.4 LTS
#>
#> Matrix products: default
#> BLAS: /usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3
#> LAPACK: /usr/lib/x86_64-linux-gnu/openblas-pthread/libopenblas-p-r0.3.26.so; LAPACK version 3.12.0
#>
#> locale:
#>  [1] LC_CTYPE=C.UTF-8      LC_NUMERIC=C          LC_TIME=C.UTF-8      LC_COLLATE=C.UTF-8   LC_
#>  [6] LC_MESSAGES=C.UTF-8  LC_PAPER=C.UTF-8     LC_NAME=C            LC_ADDRESS=C         LC_
#> [11] LC_MEASUREMENT=C.UTF-8 LC_IDENTIFICATION=C
#>
#> time zone: UTC
#> tzcode source: system (glibc)
#>
#> attached base packages:
#> [1] stats      graphics  grDevices  utils      datasets  methods    base
#>
#> other attached packages:
#>  [1] quanteda.textstats_0.97.2 visNetwork_2.1.4      ggraph_2.2.2          tidygraph_1.3.1
#>  [5] igraph_2.3.1              quanteda_4.4          pdftools_3.9.0        arrow_24.0.0
#>  [9] bibliometrix_5.4.0        RefManageR_1.4.0      bib2df_1.1.2.0        rcrossref_1.2.1
#> [13] gt_1.3.0                  tidytext_0.4.3        glue_1.8.1            openalexR_3.0.1
#> [17] lubridate_1.9.5           forcats_1.0.1         stringr_1.6.0         dplyr_1.2.1
#> [21] purrr_1.2.2              readr_2.2.0           tidyr_1.3.2           tibble_3.3.1
#> [25] ggplot2_4.0.3            tidyverse_2.0.0       bookdown_0.46         fs_2.1.0
#> [29] rmarkdown_2.31
#>
#> loaded via a namespace (and not attached):
#>  [1] bibtex_0.5.2              RColorBrewer_1.1-3     jsonlite_2.0.0        magrittr_2.0.5        fa
#>  [6] vctrs_0.7.3              memoise_2.0.1          askpass_1.2.1         base64enc_0.1-6      tin
#> [11] htmltools_0.5.9          contentanalysis_1.0.0  curl_7.1.0            janeaustenr_1.0.0    ce
#> [16] htmlwidgets_1.6.4        tokenizers_0.3.0       plyr_1.8.9            httr2_1.2.2          ca
#> [21] plotly_4.12.0            dimensionsR_0.0.3      mime_0.13             lifecycle_1.0.5      pk
#> [26] Matrix_1.7-0            R6_2.6.1              fastmap_1.2.0         shiny_1.13.0         di
#> [31] patchwork_1.3.2          shinycssloaders_1.1.0  rprojroot_2.1.1       SnowballC_0.7.1     lab
#> [36] urltools_1.7.3.1         timechange_0.4.0       polyclip_1.10-7       httr_1.4.8           cor
#> [41] here_1.0.2              bit64_4.8.0           withr_3.0.2           S7_0.2.2            ba
#> [46] viridis_0.6.5           ggforce_0.5.0         MASS_7.3-60.2         rappdirs_0.3.4      bil
#> [51] tools_4.4.1             otel_0.2.0            stopwords_2.3         zip_2.3.3            ht
#> [56] rentrez_1.2.4           promises_1.5.0        grid_4.4.1           stringdist_0.9.17    ger
```

#> [61] gtable_0.3.6	tzdb_0.5.0	rscopus_0.9.0	ca_0.71.1	da
#> [66] hms_1.1.4	xml2_1.5.2	utf8_1.2.6	ggrepel_0.9.8	pi
#> [71] nsyllable_1.0.1	vroom_1.7.1	later_1.4.8	tweenr_2.0.3	la
#> [76] bit_4.6.0	tidyselect_1.2.1	miniUI_0.1.2	knitr_1.51	gr
#> [81] crul_1.6.0	xfun_0.57	graphlayouts_1.2.3	DT_0.34.0	hun
#> [86] stringi_1.8.7	lazyeval_0.2.3	qpdf_1.4.1	yaml_2.3.12	eva
#> [91] codetools_0.2-20	httpcode_0.3.0	cli_3.6.6	xtable_1.8-8	di
#> [96] Rcpp_1.1.1-1.1	readxl_1.4.5	triebeard_0.4.1	XML_3.99-0.23	pa
#> [101] assertthat_0.2.1	pubmedR_1.0.2	viridisLite_0.4.3	scales_1.4.0	cr
#> [106] openxlsx_4.2.8.1	rlang_1.2.0	fastmatch_1.1-8		

term	tfidf
2019	
collabor	18.1
faculti	17.6
usag	16.2
measur	14.3
promot	11.8
2020	
ecosystem	17.9
open	17.7
innov	16.9
citat	15.3
journal	15.1
2021	
predatori	28.9
journal	25.5
knowledg	24.0
hold	22.0
univers	21.3
2022	
citat	23.0
preprint	22.4
model	20.7
acknowledg	20.6
academ	20.0
2023	
mention	27.5
bias	20.6
polici	20.0
journal	17.7
question	16.5