

1 Embeddings for Scientometrics

1.1 Learning objectives

After completing this chapter, you will be able to:

- Explain how word and document embeddings differ from bag-of-words representations
- Train simple word embeddings on a bibliometric corpus
- Use pre-trained embeddings to compute document similarity
- Apply UMAP or t-SNE to visualise document clusters in embedding space
- Assess when embeddings add value over simpler text representations

1.2 Setup

```
library(tidyverse)
library(openalexR)
library(quanteda)
library(word2vec)
library(uwot)
library(glue)
library(gt)

set.seed(20260509)

source(here::here("R", "api_helpers.R"))
source(here::here("R", "utils.R"))
source(here::here("R", "sci_palette.R"))
```

1.3 Conceptual background

Bag-of-words representations (??) treat each word as an independent dimension. “Citation analysis” and “bibliometric study” share no features despite being semantically related. **Word embeddings** address this by mapping words into a dense, low-dimensional vector space where semantically similar words are close together. The word2vec algorithm learns these representations by predicting words from their context (skip-gram) or context from words (CBOW) in a large training corpus.

Document embeddings extend the idea to whole documents. Simple approaches average the word vectors of a document’s words (mean pooling). More sophisticated approaches include Doc2Vec (Paragraph Vector), which learns document-specific vectors alongside word vectors, and transformer-based models like **SPECTER**, which is pre-trained on scientific paper titles and abstracts using citation-based training signals.

For scientometrics, embeddings enable several applications:

- **Semantic similarity:** Finding papers with similar meaning, even when they use different terminology.
- **Clustering:** Grouping papers by content in embedding space, complementing topic models (??).
- **Visualisation:** Projecting embeddings to 2D with UMAP or t-SNE to create interpretable “maps” of a corpus.
- **Anomaly detection:** Identifying papers that are semantically distant from their assigned field or cluster.

The trade-off is interpretability: while TF-IDF features are directly readable (the word “citation” has a specific weight), embedding dimensions are opaque. Embeddings are powerful for similarity and clustering but resist the kind of term-level interpretation that co-word analysis provides.

1.4 Worked example

1.4.1 Preparing text data

```
works <- oa_fetch(
  entity = "works",
  primary_location.source.id = "S148561398",
  from_publication_date = "2019-01-01",
  to_publication_date = "2023-12-31",
  options = list(sample = 400, seed = 42)
)

text_df <- works |>
  filter(!is.na(abstract), nchar(abstract) > 100) |>
  transmute(
    doc_id = id,
    title = display_name,
    text = paste(display_name, abstract, sep = ". "),
    year = year(publication_date),
    cited_by_count
  )

cat(glue("Documents for embedding: {nrow(text_df)}\n"))
```

```
#> Documents for embedding: 91
```

1.4.2 Training word2vec embeddings

```
clean_text <- text_df$text |>
  str_to_lower() |>
  str_replace_all("[^a-z ]", " ") |>
  str_squish()

model <- word2vec(clean_text, dim = 100, iter = 20, min_count = 3,
  type = "skip-gram", threads = 1)

cat(glue("Vocabulary size: {nrow(as.matrix(model))}\n"))
```

```
#> Vocabulary size: 1230
```

```
if ("citation" %in% rownames(as.matrix(model))) {
  nn <- predict(model, newdata = "citation", type = "nearest", top_n = 10)
  cat("Nearest neighbours of 'citation':\n")
  print(nn)
}
```

```
#> Nearest neighbours of 'citation':
#> $citation
#>      term1      term2 similarity rank
```

```
#> 1 citation compare 0.757 1
#> 2 citation meanings 0.753 2
#> 3 citation sentiment 0.739 3
#> 4 citation function 0.718 4
#> 5 citation citations 0.718 5
#> 6 citation periods 0.714 6
#> 7 citation total 0.706 7
#> 8 citation ratio 0.692 8
#> 9 citation wilcoxon 0.683 9
#> 10 citation classifications 0.681 10
```

1.4.3 Computing document embeddings via mean pooling

```
word_matrix <- as.matrix(model)

doc_embed <- function(text, wm) {
  words <- str_split(str_to_lower(str_replace_all(text, "[^a-z ]", " ")),
    "\\s+")[[1]]
  matched <- words[words %in% rownames(wm)]
  if (length(matched) == 0) return(rep(NA_real_, ncol(wm)))
  colMeans(wm[matched, , drop = FALSE])
}

doc_embeddings <- map(text_df$text, \(t) doc_embed(t, word_matrix))
embed_mat <- do.call(rbind, doc_embeddings)

valid <- complete.cases(embed_mat)
embed_mat <- embed_mat[valid, ]
text_df_valid <- text_df[valid, ]

cat(glue("Valid document embeddings: {nrow(embed_mat)}\n"))
```

```
#> Valid document embeddings: 91
```

1.4.4 Dimensionality reduction with UMAP

```
umap_result <- umap(embed_mat, n_neighbors = 15, min_dist = 0.1,
  n_components = 2, ret_model = FALSE)

text_df_valid$umap_x <- umap_result[, 1]
text_df_valid$umap_y <- umap_result[, 2]

ggplot(text_df_valid, aes(x = umap_x, y = umap_y, colour = factor(year))) +
  geom_point(alpha = 0.6, size = 1.5) +
  scale_colour_manual(values = palette_sci(n_distinct(text_df_valid$year))) +
  labs(x = "UMAP 1", y = "UMAP 2", colour = "Year") +
  theme_sci() +
  theme(axis.text = element_blank(), axis.ticks = element_blank())
```




Figure 1: UMAP projection of document embeddings, coloured by publication year.

1.4.5 Document similarity search

```
cosine_sim <- function(a, b) sum(a * b) / (sqrt(sum(a^2)) * sqrt(sum(b^2)))

query_idx <- 1
sims <- map_dbl(seq_len(nrow(embed_mat)),
               \(i) cosine_sim(embed_mat[query_idx, ], embed_mat[i, ]))

top_similar <- order(sims, decreasing = TRUE)[2:6]

cat(glue("Query: {text_df_valid$title[query_idx]}\n\n"))
```

```
#> Query: Towards medical knowmetrics: representing and computing medical knowledge using semantic pred
```

```
cat("Most similar documents:\n")
```

```
#> Most similar documents:
```

```
walk(top_similar, \(i) {
  cat(glue("  [{round(sims[i], 3)}] {text_df_valid$title[i]}\n"))
})
```

```
#> [0.9] Draw me Science[0.896] Linking the dimensions of policy-related research on obesity: a hybrid r
```

1.4.6 Embedding clusters vs. citation impact

```
ggplot(text_df_valid, aes(x = umap_x, y = umap_y,
                          colour = log1p(cited_by_count))) +
  geom_point(alpha = 0.6, size = 1.5) +
  scale_colour_viridis_c(option = "C", name = "log(cites + 1)") +
  labs(x = "UMAP 1", y = "UMAP 2") +
  theme_sci() +
  theme(axis.text = element_blank(), axis.ticks = element_blank())
```

1.5 Diagnostics and interpretation

- **Vocabulary coverage:** Check what fraction of abstract words appear in the embedding vocabulary. Low coverage (< 80%) indicates the training corpus is too small or preprocessing is too aggressive.
- **Nearest neighbours:** Inspect the nearest neighbours of key domain terms. If “citation” is closest to “references” and “bibliometric”, the embeddings capture domain structure.
- **UMAP stability:** UMAP is stochastic. Run with multiple seeds and check whether the global cluster structure is stable. Individual point positions will vary but clusters should persist.
- **Embedding dimensionality:** 100–300 dimensions is standard. Below 50, embeddings lose semantic resolution; above 300, training becomes slow with diminishing returns.




Figure 2: UMAP embedding coloured by citation impact (log scale).

1.6 Limitations and responsible use

1.7 Limitations and responsible use

- **Embeddings are opaque.** Unlike TF-IDF weights, embedding dimensions have no interpretable meaning. You can measure similarity but not explain *why* two documents are similar without additional analysis.
- **Training data bias.** Embeddings reflect the biases of their training corpus. If the corpus overrepresents certain topics or perspectives, the embeddings will too.
- **Small corpora produce poor embeddings.** Word2vec needs tens of thousands of documents for reliable training. For small bibliometric corpora (< 1,000 papers), use pre-trained embeddings instead.
- **UMAP/t-SNE distort distances.** These methods preserve local structure but distort global distances. Documents that appear close in 2D may not be the most similar in high-dimensional space [hicks2015].

1.8 Common pitfalls

1.9 Common pitfalls

- **Training on too few documents.** Word2vec on 200 abstracts produces unreliable embeddings. Use at least 5,000 documents or switch to pre-trained models.
- **Averaging without weighting.** Simple mean pooling treats every word equally. TF-IDF-weighted averaging gives more weight to distinctive terms and often produces better document embeddings.
- **Over-interpreting UMAP clusters.** Visual clusters in UMAP may reflect projection artifacts. Validate clusters with a separate method (e.g., k-means in the full embedding space).
- **Comparing embeddings from different models.** Embedding spaces are not aligned across different training runs. Cosine similarity is only meaningful within a single embedding space.

1.10 Exercises

1. **Weighted averaging.** Compute document embeddings using TF-IDF-weighted word vectors instead of simple means. Does the UMAP visualisation change?
2. **K-means clustering.** Apply k-means ($k = 5, 8, 12$) to the document embeddings. Compare the resulting clusters with LDA topics from ???. Do they agree?
3. **Pre-trained embeddings.** If available, use SPECTER embeddings (via reticulate and Python) instead of corpus-trained word2vec. How does similarity search quality change?
4. **Analogy test.** Test whether the word2vec model captures analogical relationships (e.g., “journal” - “article” + “book” “publisher”). What does the result tell you about the model?

1.11 Solutions

Solutions are provided in 1.11.

1.12 Further reading

- @silge2017 — Text embeddings in R, including integration with word2vec.
- @priem2022 — OpenAlex concept embeddings and their relationship to document content.
- @waltman2012 — Bibliometric similarity measures as context for embedding-based approaches.

1.13 Session info

```
sessionInfo()
```

```
#> R version 4.4.1 (2024-06-14)
#> Platform: x86_64-pc-linux-gnu
#> Running under: Ubuntu 24.04.4 LTS
#>
#> Matrix products: default
#> BLAS: /usr/lib/x86_64-linux-gnu/openblas-pthread/libblas.so.3
#> LAPACK: /usr/lib/x86_64-linux-gnu/openblas-pthread/libopenblas-p-r0.3.26.so; LAPACK version 3.12.0
#>
#> locale:
#>  [1] LC_CTYPE=C.UTF-8          LC_NUMERIC=C              LC_TIME=C.UTF-8
#>  [4] LC_COLLATE=C.UTF-8       LC_MONETARY=C.UTF-8      LC_MESSAGES=C.UTF-8
#>  [7] LC_PAPER=C.UTF-8         LC_NAME=C                 LC_ADDRESS=C
#> [10] LC_TELEPHONE=C           LC_MEASUREMENT=C.UTF-8   LC_IDENTIFICATION=C
#>
#> time zone: UTC
#> tzcode source: system (glibc)
#>
#> attached base packages:
#> [1] stats      graphics  grDevices  utils      datasets  methods    base
#>
#> other attached packages:
#>  [1] uwot_0.2.4             Matrix_1.7-0
#>  [3] word2vec_0.4.1         stm_1.3.8
#>  [5] topicmodels_0.2-17     quanteda.textstats_0.97.2
#>  [7] visNetwork_2.1.4       ggraph_2.2.2
#>  [9] tidygraph_1.3.1        igraph_2.3.2
#> [11] quanteda_4.4           pdftools_3.9.0
#> [13] arrow_24.0.0           bibliometrix_5.4.0
#> [15] RefManager_1.4.0       bib2df_1.1.2.0
#> [17] rcrossref_1.2.1        gt_1.3.0
#> [19] tidytext_0.4.3         glue_1.8.1
#> [21] openalexR_3.0.1        lubridate_1.9.5
#> [23] forcats_1.0.1          stringr_1.6.0
#> [25] dplyr_1.2.1            purrr_1.2.2
#> [27] readr_2.2.0            tidyr_1.3.2
#> [29] tibble_3.3.1           ggplot2_4.0.3
#> [31] tidyverse_2.0.0        bookdown_0.46
#> [33] fs_2.1.0               rmarkdown_2.31
#>
#> loaded via a namespace (and not attached):
#>  [1] bibtex_0.5.2           RColorBrewer_1.1-3      jsonlite_2.0.0
#>  [4] magrittr_2.0.5         modeltools_0.2-24       farver_2.1.2
#>  [7] vctr_0.7.3            memoise_2.0.1           askpass_1.2.1
#> [10] base64enc_0.1-6        tinytex_0.59            htmltools_0.5.9
#> [13] contentanalysis_1.0.0  curl_7.1.0              janeaustenr_1.0.0
#> [16] cellranger_1.1.0       htmlwidgets_1.6.4       tokenizers_0.3.0
#> [19] plyr_1.8.9            http2_1.2.2             cachem_1.1.0
#> [22] plotly_4.12.0          dimensionsR_0.0.3       mime_0.13
#> [25] lifecycle_1.0.5        pkgconfig_2.0.3         R6_2.6.1
```

#> [28] fastmap_1.2.0	shiny_1.13.0	digest_0.6.39
#> [31] patchwork_1.3.2	shinycssloaders_1.1.0	rprojroot_2.1.1
#> [34] RSpectra_0.16-2	SnowballC_0.7.1	labeling_0.4.3
#> [37] urltools_1.7.3.1	timechange_0.4.0	polyclip_1.10-7
#> [40] httr_1.4.8	compiler_4.4.1	here_1.0.2
#> [43] bit64_4.8.0	withr_3.0.2	S7_0.2.2
#> [46] backports_1.5.1	viridis_0.6.5	ggforce_0.5.0
#> [49] MASS_7.3-60.2	rappdirs_0.3.4	bibliometrixData_0.3.0
#> [52] tools_4.4.1	otel_0.2.0	stopwords_2.3
#> [55] zip_2.3.3	httpuv_1.6.17	rentrez_1.2.4
#> [58] promises_1.5.0	grid_4.4.1	stringdist_0.9.17
#> [61] reshape2_1.4.5	generics_0.1.4	gtable_0.3.6
#> [64] tzdb_0.5.0	rscopus_0.9.0	ca_0.71.1
#> [67] data.table_1.18.4	hms_1.1.4	xml2_1.5.2
#> [70] utf8_1.2.6	ggrepel_0.9.8	pillar_1.11.1
#> [73] nsyllable_1.0.1	vroom_1.7.1	later_1.4.8
#> [76] tweenr_2.0.3	lattice_0.22-6	FNN_1.1.4.1
#> [79] bit_4.6.0	tidyselect_1.2.1	tm_0.7-18
#> [82] miniUI_0.1.2	knitr_1.51	gridExtra_2.3
#> [85] NLP_0.3-2	stats4_4.4.1	crul_1.6.0
#> [88] xfun_0.57	graphlayouts_1.2.3	matrixStats_1.5.0
#> [91] DT_0.34.0	humaniformat_0.6.0	stringi_1.8.7
#> [94] lazyeval_0.2.3	qpdf_1.4.1	yaml_2.3.12
#> [97] evaluate_1.0.5	codetools_0.2-20	httpcode_0.3.0
#> [100] cli_3.6.6	xtable_1.8-8	dichromat_2.0-0.1
#> [103] Rcpp_1.1.1-1.1	readxl_1.4.5	triebeard_0.4.1
#> [106] XML_3.99-0.23	parallel_4.4.1	assertthat_0.2.1
#> [109] pubmedR_1.0.2	slam_0.1-55	viridisLite_0.4.3
#> [112] scales_1.4.0	crayon_1.5.3	openxlsx_4.2.8.1
#> [115] rlang_1.2.0	fastmatch_1.1-8	